


E-Commerce Churn Analysis

Team Project: Beta

1. Andhika Prakoso
2. Emil Supriatna
3. Monica Kristin Napitupulu



Content of Presentation:



Business Problem

Apa Masalah yang dihadapi?
Bagaimana menyelesaikannya?

Data Preprocessing

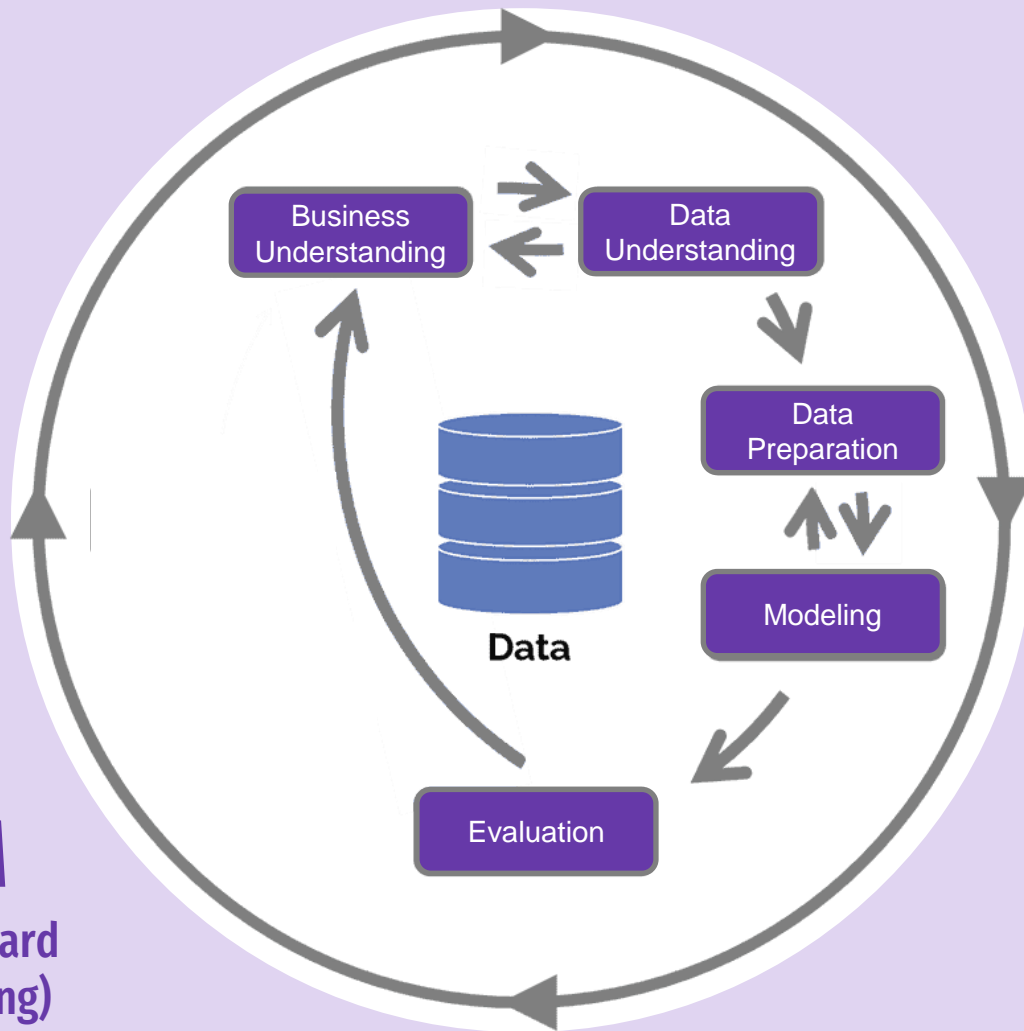
Bagaimana data dipahami?
Bagaimana data disiapkan untuk
analisis tingkat lanjut?

Modeling & Evaluation

Apa model yang tepat digunakan?
Bagaimana kontribusinya terhadap
permasalahan?

Recommendation

Apa rekomendasi bisnis yang
dapat diberikan pada Stakeholder?



CRISP-DM

(Cross-Industry Standard
Process for Data Mining)

01

Business Problem

Apa Masalah yang dihadapi?
Bagaimana menyelesaikannya?

*"Business understanding to
understand the problem"*



Predict Churn to Improve Business Revenue

Business revenue stream

→ Customer Acquisition

→ Customer Retention

Manfaat Customer Retention

5x biaya lebih rendah

Meningkatkan profit 25% - 95%

Kemudahan cross-selling produk

Source: <https://hbr.org/2014/10/the-value-of-keeping-the-right-customers>



Business Problem:

- Fenomena Churn di dalam bisnis tidak dapat dihindari. Namun angka Churn bisa ditekan melalui berbagai macam strategi.
- Agar dapat memberikan strategi yang efektif, **bisnis membutuhkan suatu sistem prediksi yang mampu membantu untuk mengurangi tingkat churn.**


Goals and Analytics Approach

Model Prediction



Membangun model
prediksi terhadap potensi
Churn di masa depan

Influential Factors



Mengetahui faktor-faktor
paling berpengaruh
terhadap Churn

Matrics Evaluation

Recall

Recall menyatakan seberapa besar persentase kejadian pada kelas positif yang berhasil dideteksi.


Cara kerjanya adalah semakin tinggi skor Recall, semakin kecil nilai False Negative, yang berarti model semakin baik dalam mendeteksi kasus Churn secara tepat.

Rumus:

$$Recall = \frac{TP}{TP + FN}$$

		PREDICTED	
		Negative	Positive
ACTUAL	Negative	True Negative (TN)	False Positive (FP) Type I Error
	Positive	False Negative (FN) Type II Error	True Positive (TP)

Confusion Matrix



02

Data Preprocessing

Step 1: Data Understanding
Step 2: Data Preparation

Step 1:
Data Understanding

Data Understanding

	Variable	Description
0	CustomerID	Unique customer ID
1	Churn	Churn Flag
2	Tenure	Tenure of customer in organization
3	PreferredLoginDevice	Preferred login device of customer
4	CityTier	City tier
5	WarehouseToHome	Distance in between warehouse to home of customer
6	PreferredPaymentMode	Preferred payment method of customer
7	Gender	Gender of customer
8	HourSpendOnApp	Number of hours spend on mobile application or website
9	NumberOfDeviceRegistered	Total number of deceives is registered on particular customer
10	PreferedOrderCat	Preferred order category of customer in last month
11	SatisfactionScore	Satisfactory score of customer on service
12	MaritalStatus	Marital status of customer
13	NumberOfAddress	Total number of added added on particular customer
14	Complain	Any complaint has been raised in last month
15	OrderAmountHikeFromlastYear	Percentage increases in order from last year
16	CouponUsed	Total number of coupon has been used in last month
17	OrderCount	Total number of orders has been places in last month
18	DaySinceLastOrder	Day Since last order by customer
19	CashbackAmount	Average cashback in last month

Data Type Checking

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 5630 entries, 0 to 5629
```

```
Data columns (total 20 columns):
```

#	Column	Non-Null Count	Dtype
0	CustomerID	5630 non-null	int64
1	Churn	5630 non-null	int64
2	Tenure	5366 non-null	float64
3	PreferredLoginDevice	5630 non-null	object
4	CityTier	5630 non-null	int64
5	WarehouseToHome	5379 non-null	float64
6	PreferredPaymentMode	5630 non-null	object
7	Gender	5630 non-null	object
8	HourSpendOnApp	5375 non-null	float64
9	NumberOfDeviceRegistered	5630 non-null	int64
10	PreferedOrderCat	5630 non-null	object
11	SatisfactionScore	5630 non-null	int64
12	MaritalStatus	5630 non-null	object
13	NumberOfAddress	5630 non-null	int64
14	Complain	5630 non-null	int64
15	OrderAmountHikeFromlastYear	5365 non-null	float64
16	CouponUsed	5374 non-null	float64
17	OrderCount	5372 non-null	float64
18	DaySinceLastOrder	5323 non-null	float64
19	CashbackAmount	5630 non-null	float64

```
dtypes: float64(8), int64(7), object(5)
```













```
memory usage: 879.8+ KB
```

	count	mean	std	min	25%	50%	75%	max
Tenure	5366.0	10.189899	8.557241	0.0	2.00	9.00	16.0000	61.00
CityTier	5630.0	1.654707	0.915389	1.0	1.00	1.00	3.0000	3.00
WarehouseToHome	5379.0	15.639896	8.531475	5.0	9.00	14.00	20.0000	127.00
HourSpendOnApp	5375.0	2.931535	0.721926	0.0	2.00	3.00	3.0000	5.00
NumberOfDeviceRegistered	5630.0	3.688988	1.023999	1.0	3.00	4.00	4.0000	6.00
SatisfactionScore	5630.0	3.066785	1.380194	1.0	2.00	3.00	4.0000	5.00
NumberOfAddress	5630.0	4.214032	2.583586	1.0	2.00	3.00	6.0000	22.00
Complain	5630.0	0.284902	0.451408	0.0	0.00	0.00	1.0000	1.00
OrderAmountHikeFromLastYear	5365.0	15.707922	3.675485	11.0	13.00	15.00	18.0000	26.00
CouponUsed	5374.0	1.751023	1.894621	0.0	1.00	1.00	2.0000	16.00
OrderCount	5372.0	3.008004	2.939680	1.0	1.00	2.00	3.0000	16.00
DaySinceLastOrder	5323.0	4.543491	3.654433	0.0	2.00	3.00	7.0000	46.00
CashbackAmount	5630.0	177.223030	49.207036	0.0	145.77	163.28	196.3925	324.99

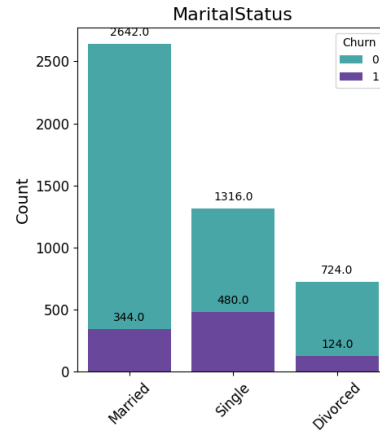

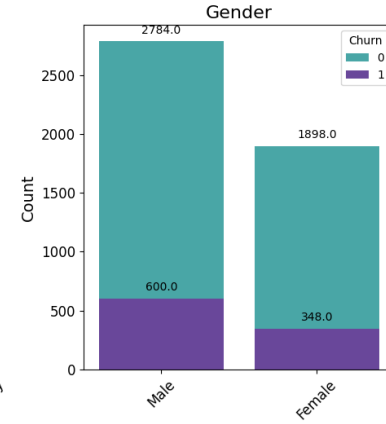


Simple Statistic

	count	unique	top	freq
PreferredLoginDevice	5630	3	Mobile Phone	2765
PreferredPaymentMode	5630	7	Debit Card	2314
Gender	5630	2	Male	3384
PreferedOrderCat	5630	6	Laptop & Accessory	2050
MaritalStatus	5630	3	Married	2986


Data Distribution (numeric)



Data Distribution (categoric)










Data Correlation

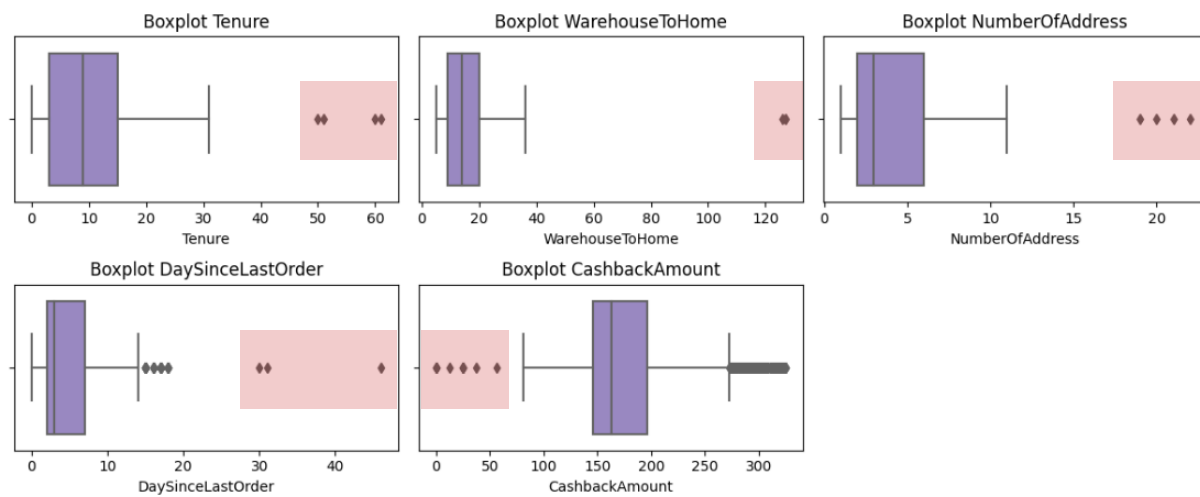


Step 2:
Data Preparation

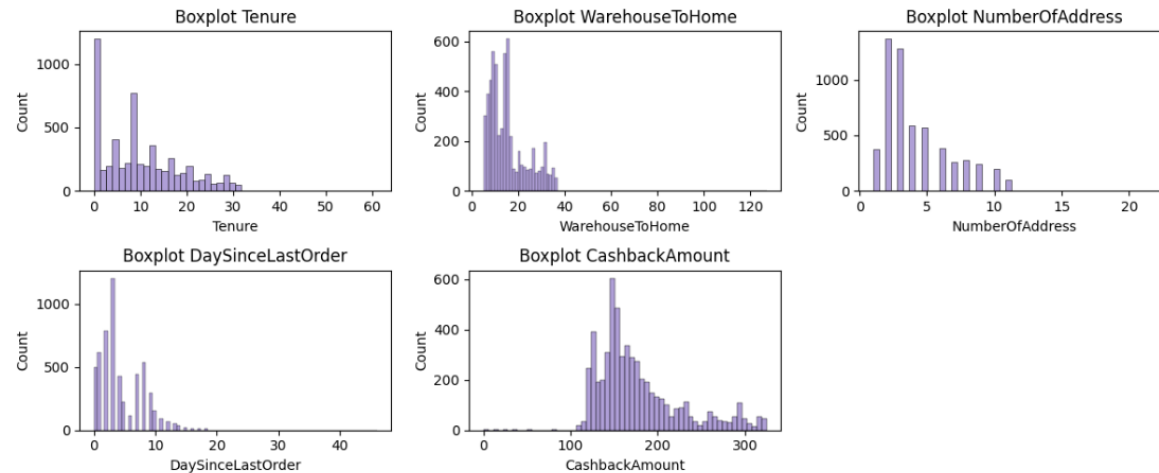
Missing Values Handling

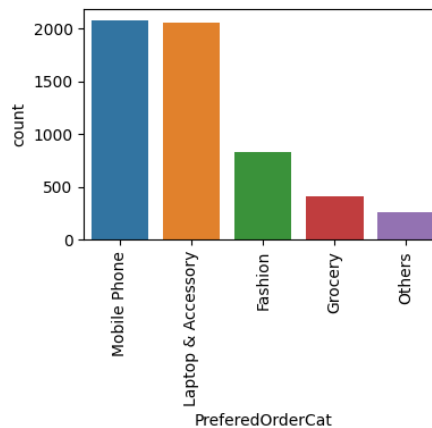







```
# Inisialisasi SimpleImputer  
imputer = SimpleImputer(strategy='median')
```

Outliers Handling





Inconsistent Variable

Other Data Preparation

DATA DUPLIKAT

```
duplicated_data = df.duplicated(subset='CustomerID')  
print(df[duplicated_data])
```

Empty DataFrame

DROP CUSTOMER ID


```
# Drop Customer ID  
df = df.drop(['CustomerID'], axis=1)
```

MAPPING FOR VISUALIZATION

```
df['CityTier'] = df['CityTier'].replace({1: 'Tier_1', 2: 'Tier_2', 3: 'Tier_3'})  
df['SatisfactionScore'] = df['SatisfactionScore'].replace({1: 'Very_Satisfied',  
2: 'Satisfied',  
3: 'Neutral',  
4: 'Dissatisfied',  
5: 'Highly_Dissatisfied'})  
df['Complain'] = df['Complain'].replace({1: 'Complain', 0: 'non-Complain'})  
df['PreferredOrderCat'] = df['PreferredOrderCat'].replace('Laptop & Accessory', 'Laptop & Acc')  
df['NumberOfDeviceRegistered'] = df['NumberOfDeviceRegistered'].replace({1: '1_Device',  
2: '2_Devices',  
3: '3_Devices',  
4: '4_Devices',  
5: '5_Devices',  
6: '6_Devices'})
```

EDA for Modelling

Train-Val Comparison to Data Distribution




ENCODING DAN SCALING

```
ordinal_mapping = [
    {"col": "CityTier", "mapping": {"Tier_1": 1, "Tier_2": 2, "Tier_3": 3}},
    {"col": "NumberOfDeviceRegistered", "mapping": {"1-2_Devices": 1, "3-4_Devices": 2, "5_Devices": 3, "6_Devices": 4}},
    {"col": "SatisfactionScore", "mapping": {"Satisfied & more_than": 1, "Neutral": 2, "Dissatisfied": 3,
                                             "Highly_Dissatisfied": 4}},
    {"col": "PreferedOrderCat", "mapping": {"Grocery": 1, "Fashion": 2, "Mobile Phone": 3, "Laptop, ACC & Others": 4}}]

transformer = ColumnTransformer([
    ('One Hot Encoding', OneHotEncoder(drop='first'), ['PreferredLoginDevice', 'PreferredPaymentMode',
                                                         'MaritalStatus', 'Complain']),
    ('Ordinal Encoding', ce.OrdinalEncoder(cols=['CityTier', 'NumberOfDeviceRegistered',
                                                  'SatisfactionScore', 'PreferedOrderCat'],
                                          mapping=ordinal_mapping), ['CityTier', 'NumberOfDeviceRegistered',
                                                                      'SatisfactionScore', 'PreferedOrderCat']),
    ('Robust', RobustScaler(), ['Tenure', 'DaySinceLastOrder', 'CashbackAmount']),
])
```

ENCODED and SCALED FEATURES

#	Column	Non-Null Count	Dtype
0	One Hot Encoding__PreferredLoginDevice_Mobile Phone	3923 non-null	float64
1	One Hot Encoding__PreferredPaymentMode_Credit-C & UPI	3923 non-null	float64
2	One Hot Encoding__PreferredPaymentMode_Debit-C & eWallet	3923 non-null	float64
3	One Hot Encoding__MaritalStatus_Married	3923 non-null	float64
4	One Hot Encoding__MaritalStatus_Single	3923 non-null	float64
5	One Hot Encoding__Complain_non-Complain	3923 non-null	float64
6	Ordinal Encoding__CityTier	3923 non-null	float64
7	Ordinal Encoding__NumberOfDeviceRegistered	3923 non-null	float64
8	Ordinal Encoding__SatisfactionScore	3923 non-null	float64
9	Ordinal Encoding__PreferedOrderCat	3923 non-null	float64
10	Robust__Tenure	3923 non-null	float64
11	Robust__DaySinceLastOrder	3923 non-null	float64
12	Robust__CashbackAmount	3923 non-null	float64



03

Modeling & Evaluation

Apa model yang tepat digunakan?
Bagaimana evaluasi dan kontribusinya
terhadap permasalahan?

Model Benchmarking: K-Fold

	mean recall	StdDev
model		
XGBoost	0.794	0.059
AdaBoost	0.779	0.031
Decision Tree	0.778	0.045
Random Forest	0.754	0.064
LightGBM	0.749	0.059
GBoost	0.572	0.070
KNN	0.466	0.025

Model Benchmarking: Data Validation

	Model	recall score
1	Decision Tree	0.838
5	XGBoost	0.824
3	AdaBoost	0.817
2	Random Forest	0.782
6	LightGBM	0.754
4	GBoost	0.542
0	KNN	0.521

Oversampling

```
# Import oversampling method  
Smote = SMOTE(random_state = 2020)  
Ros = RandomOverSampler(random_state=2020)
```

	mean recall Score	StdDev
model		
XGB_ros	0.838	0.063
XGB	0.794	0.059
XGB_smote	0.784	0.036

	mean recall Score	StdDev
model		
TREE_smote	0.781	0.028
TREE	0.778	0.045
TREE_ros	0.758	0.050

	mean recall Score	StdDev
model		
ADA_smote	0.781	0.032
ADABOOST	0.779	0.031
ADA_ros	0.763	0.055

feature selection

```
feature_selection_XGB = RFE(estimator=XGBClassifier(random_state=2020), n_features_to_select=11)
```

	mean recall Score	StdDev
model		
XGB_ROS_FS	0.843	0.060
XGB_ROS	0.838	0.063

	mean recall Score	StdDev
model		
TREE_SMOTE_FS	0.775	0.057
TREE_SMOTE	0.758	0.050

	mean recall Score	StdDev
model		
ADA_FS	0.785	0.043
model	0.779	0.031

Machine Learning we use



XGBoost is an optimized distributed gradient boosting library designed for efficient and scalable training of machine learning models. It is an ensemble learning method that combines the predictions of multiple weak models to produce a stronger prediction.
(geeksforgeeks.org/xgboost)

```
# Hyperparameter Tuning pada XGB dengan Random Over Sampling dan Feature Selection
xgb = XGBClassifier(random_state=2020)
```

```
estimator = Pipeline([
    ('preprocess', transformer),
    ('resampler', Ros),
    ('feature_selector', feature_selection_XGB),
    ('model', xgb)])
```


```
hyperparam_space = {
    'model__n_estimators': [100, 200, 300],
    'model__max_depth': [3, 5, 7],
    'model__learning_rate': [0.01, 0.1, 0.3],
    'model__subsample': [0.7, 0.8, 0.9],
}
```

```
grid_search_xgb = GridSearchCV(
    estimator,
    param_grid = hyperparam_space,
    cv = skfold,
    scoring = 'recall',
    n_jobs = -1
)
```

Recall Score XGB Default dengan Feature Selection : 0.8873239436619719
Recall Score XGB Tuned dengan Feature Selection: 0.8591549295774648


Recall Score XGB Default tanpa feature selection: 0.9154929577464789
Recall Score XGB Tuned tanpa feature selection: 0.8873239436619719

Final Report




Classification Report Default XGB tanpa feature selection :					
	precision	recall	f1-score	support	
0	0.98	0.97	0.98	699	
1	0.86	0.92	0.88	142	
accuracy			0.96	841	
macro avg	0.92	0.94	0.93	841	
weighted avg	0.96	0.96	0.96	841	

Feature Importance



Feature
Importance

SHAP
Feature




Feature Limitation

	min	max
Tenure	0.00	31.00
DaySinceLastOrder	0.00	18.00
CashbackAmount	110.09	324.99

	dataFeatures	uniqueSample
0	PreferredLoginDevice	[Mobile Phone, Computer]
1	CityTier	[Tier_1, Tier_3, Tier_2]
2	PreferredPaymentMode	[Debit-C & eWallet, Credit-C & UPI, COD]
3	NumberOfDeviceRegistered	[3-4_Devices, 5_Devices, 1-2_Devices, 6_Devices]
4	PreferredOrderCat	[Grocery, Fashion, Laptop, ACC & Others, Mobile Phone]
5	SatisfactionScore	[Neutral, Dissatisfied, Satisfied & more_than, Highly_Dissatisfied]
6	MaritalStatus	[Single, Married, Divorced]
7	Complain	[non-Complain, Complain]

Rule-Based vs Model



Churn Cost
40%
Increase




Model can predict
2.3x
More Churn than
Rule-Based

Rule-Based and Model Simulation

Based on Data Test (841 data)

Rule-Based Simulation



```
# Cek logika dari data nilai negatif menggunakan csv
rule_based = data_test[(data_test['Tenure'] < 7.75) &
                        (data_test['Complain'] == 'Complain')]

print(f'Metode Pemberian Cashback pada Customer Churn dengan Metode RULE-BASED:\n')
print(f'Rule-Based Churn Cost: Rp{(rule_based.shape[0]*budget_perCust):,} atau sekitar {rule_based.shape[0]} Customer.')
print(f'Jumlah customer yang benar-benar churn dan tercover cashback: {rule_based["Churn"].value_counts()[1]} Customer')
print(f'Biaya yang terbuang karena salah memperkirakan customer churn: Rp{(rule_based.shape[0]-rule_based["Churn"].value_counts()[0]):,}
```


Metode Pemberian Cashback pada Customer Churn dengan Metode RULE-BASED:

Rule-Based Churn Cost: Rp5,400,000 atau sekitar 108 Customer.

Jumlah customer yang benar-benar churn dan tercover cashback: 56 Customer

Biaya yang terbuang karena salah memperkirakan customer churn: Rp2,600,000

Model Simulation



```
tn, fp, fn, tp = confusion_matrix(y_test, y_pred_default_xgb2).ravel()
```

```
print(f'Metode Pemberian Cashback pada Customer Churn dengan MODELLING:\n')
print(f'Modelling Churn Cost: Rp{(budget_perCust*tp)+(budget_perCust*fp):,}')
print(f'Jumlah customer yang benar-benar churn dan tercover cashback: {tp} (recall 1 = {tp/(fn+tp):.2f})')
print(f'Total customer churn yang gagal diprediksi oleh model: {fn} (false negatif)')
print(f'Biaya yang terbuang Rp{budget_perCust*fp:,} karena salah memprediksi sebanyak {fp} customer (false positif)')
```


Metode Pemberian Cashback pada Customer Churn dengan MODELLING:

Modelling Churn Cost: Rp7,600,000

Jumlah customer yang benar-benar churn dan tercover cashback: 130 (recall 1 = 0.92)

Total customer churn yang gagal diprediksi oleh model: 12 (false negatif)

Biaya yang terbuang Rp1,100,000 karena salah memprediksi sebanyak 22 customer (false positif)



04

Recommendation

Apa rekomendasi bisnis yang dapat diberikan pada Stakeholder?

Business Recommendation



Business Recommendation




Model Recommendation

01

Performa model mungkin dapat diperbaiki melalui pengujian kembali *hyperparam_space*

02

Menambahkan beberapa fitur lain untuk meningkatkan akurasi model seperti *Last_Login*, *Total_purchase*, *Total_product_type_purchased*



04

Perbaiki performa model mungkin dapat coba dilakukan melalui *tuning* pada 2 model teratas lainnya (*Decision Tree* dan *Adabost*)

03

Perbanyak data untuk kualitas hasil modelling yang lebih baik. Termasuk perbaiki kualitas data seperti mengurangi *missing values* dan *error labels*.



Thank you!